RULLMR Session 05: Data analysis with LLMs

Preparation for 2 November 2026

This week we pretend we're dealing with **sensitive** data we want to analyse and report in a **paper** (so keep e.g. privacy and reproducibility in mind).

As homework, because there's not a lot of time until the next session, its enough if you sketch out **a workflow/steps to solve the task**. You can also do the whole task if you're eager, but the goal of the next session will be mostly to do it in teams. The data can be found here under session 05: https://bma-vandijk.github.io/RULLMR/schedule/.

We're not giving extensive instructions or tips because we want to compare your own approaches and considerations during class. The dataset are simple because you're a broad and mixed audience; the process is as important as the outcome. Choose one task that is closest to your interests.

Task 1 – Prediction model with tabular data
We want to predict length of stay since admission with a sensitive ED dataset called hospital_admissions.csv and a linear model.

Task 2 – Text classification
We want to predict smoking status from a set of clinical notes called clinical_notes.csv. Smoking means smoking or vaping anything, from cigars to roll-ups and e-cigarettes.

Every text note should get one of the following labels:
CURRENT: smokes now – even occasionally
FORMER: used to smoke but has stopped (no matter how long ago)
NEVER: has never smoked or vaped
UNKNOWN: note does not say anything about patient's smoking

Note: gold_labels.csv contains the answer key that you can use to verify your approach.